Multimedia Systems

Dr Laura Toni
l.toni@ucl.ac.uk

University College London

Today Lecture

Introduction to Information compression

- Source Coding
- Information and Entropy
- Variable length coding
- Quantization

Multimedia Systems

- Image and Lossy Compression
  - Transforms
  - JPEG Quantization
  - JPEG Lossless Compression

- Video Compression
  - Motion Compensation

Block Diagram for Image Compression

Signal Decomposition:

- Compact energy into a small number of coefficients
- Decorrelate the components of the signal

Quantization:

- Make approximations to the transform coefficients
- Selectively throw away information

Lossless coding

- Variable length coding (e.g., Huffman Coding)
- Choose short/long codewords to minimize number of bits

Transform Coding

- We are interested in transforming a signal or image from the **spatial domain** to the **frequency domain**
- by itself, the transform – does not save any bits
  – does not introduce any distortion
- both of these happen when we throw away information - “lossy compression” implemented by the quantizer

2-d DCT Basis Functions

$$C(u,v)=\alpha(u)\alpha(v)\sum_{x=0}^{N-1}\sum_{y=0}^{N-1} f(x,y)\cos\left[\frac{(2x+1)u\pi}{2N}\right]\cos\left[\frac{(2y+1)v\pi}{2N}\right]$$

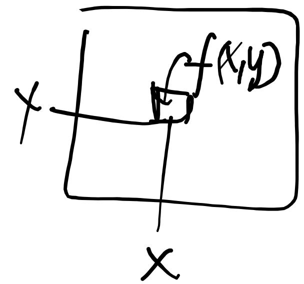

- The whole array of basis functions looks like this:

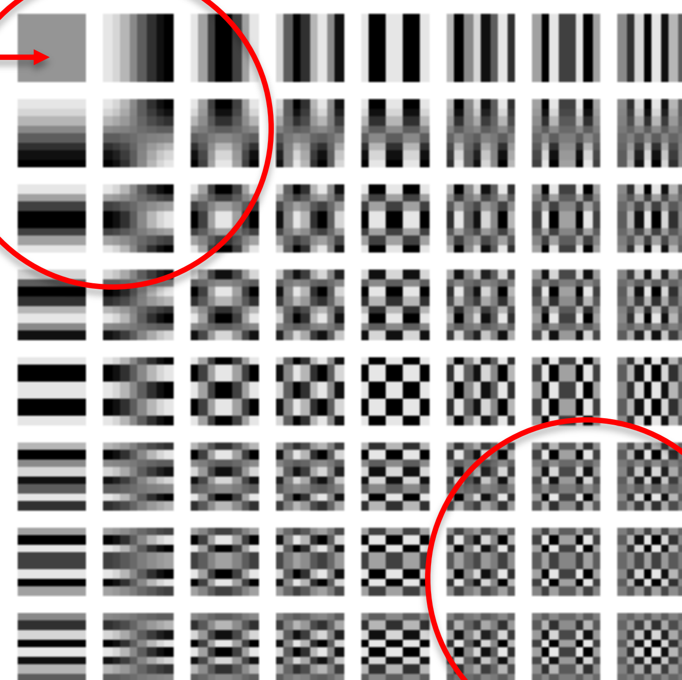

The array of DCT coefficients tells how much of each of these basis images we need to regenerate original image

$$x_t \to \boxed{\mathcal{F}} \to \underline{X(f)} \to \boxed{\mathcal{F}^{-1}} \to x_t$$

Why is the DCT a good idea?

smooth

edge

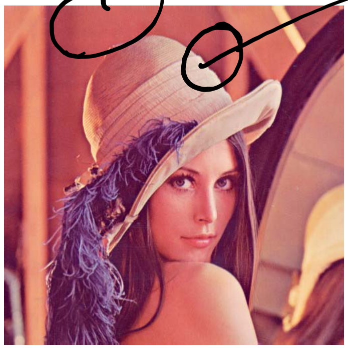

- Smooth areas in the images have pixels highly correlated
- Pixels around edges are less correlated

- Our visual system is less sensitive to distortion around edges
- We want to protect smooth areas the most

- DCT has excellent energy compaction for highly correlated data (smooth areas).
- Low frequencies in the DCT correspond to smooth areas
- We can quantize less heavily low frequencies

low fr. ⟵ smooth ⟶ more sensitive

hig fr. ⟵ edge ⟶ less sensitive

Quantization after Transforming

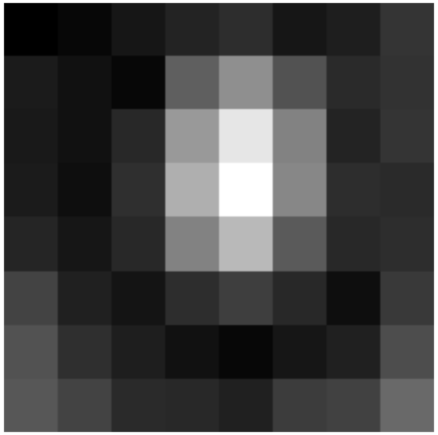

$$\begin{array}{cc} & N \\ N & \begin{bmatrix} 52 & 55 & 61 & 66 & 70 & 61 & 64 & 73 \\ 63 & 59 & 55 & 90 & 109 & 85 & 69 & 72 \\ 62 & 59 & 68 & 113 & 144 & 104 & 66 & 73 \\ 63 & 58 & 71 & 122 & 154 & 106 & 70 & 69 \\ 67 & 61 & 68 & 104 & 126 & 88 & 68 & 70 \\ 79 & 65 & 60 & 70 & 77 & 68 & 58 & 75 \\ 85 & 71 & 64 & 59 & 55 & 61 & 65 & 83 \\ 87 & 79 & 69 & 68 & 65 & 76 & 78 & 94 \end{bmatrix} \end{array}$$

DCT

$$\begin{array}{cc} & N \\ N & \begin{bmatrix} -415 & -30 & -61 & 27 & 56 & -20 & -2 & 0 \\ 4 & -22 & -61 & 10 & 13 & -7 & -9 & 5 \\ -47 & 7 & 77 & -25 & -29 & 10 & 5 & -6 \\ -49 & 12 & 34 & -15 & -10 & 6 & 2 & 2 \\ 12 & -7 & -13 & -4 & -2 & 2 & -3 & 3 \\ -8 & 3 & 2 & -6 & -2 & 1 & 4 & 2 \\ -1 & 0 & 0 & -2 & -1 & -3 & 4 & -1 \\ 0 & 0 & -1 & -4 & -1 & 0 & 1 & 2 \end{bmatrix} \end{array}$$

- After the DCT, we still have N^2 numbers to represent a block:

  N^2 pixels $\leftrightarrow$ N^2 coefficients

- However, the N^2 coefficients will have
  - Lot of values near zero
  - Lot of values which represent high frequency info

Block Diagram for Image Compression

Original Image → Signal Decomposition → Quantization → Lossless Coding → Compressed Bit Stream

Signal Decomposition:

- Compact energy into a small number of coefficients
- Decorrelate the components of the signal

Quantization:

- Make approximations to the transform coefficients
- Selectively throw away information

Lossless coding

- Variable length coding (e.g., Huffman Coding)
- Choose short/long codewords to minimize number of bits

Let's continue the example

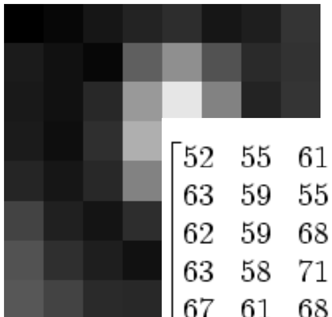

$$\begin{bmatrix} 52 & 55 & 61 & 66 & 70 & 61 & 64 & 73 \\ 63 & 59 & 55 & 90 & 109 & 85 & 69 & 72 \\ 62 & 59 & 68 & 113 & 144 & 104 & 66 & 73 \\ 63 & 58 & 71 & 122 & 154 & 106 & 70 & 69 \\ 67 & 61 & 68 & 104 & 126 & 88 & 68 & 70 \\ 79 & 65 & 60 & 70 & 77 & 68 & 58 & 75 \\ 85 & 71 & 64 & 59 & 55 & 61 & 65 & 83 \\ 87 & 79 & 69 & 68 & 65 & 76 & 78 & 94 \end{bmatrix} \xrightarrow{} \boxed{\text{DCT}} \xrightarrow{} \begin{bmatrix} -415 & -30 & -61 & 27 & 56 & -20 & -2 & 0 \\ 4 & -22 & -61 & 10 & 13 & -7 & -9 & 5 \\ -47 & 7 & 77 & -25 & -29 & 10 & 5 & -6 \\ -49 & 12 & 34 & -15 & -10 & 6 & 2 & 2 \\ 12 & -7 & -13 & -4 & -2 & 2 & -3 & 3 \\ -8 & 3 & 2 & -6 & -2 & 1 & 4 & 2 \\ -1 & 0 & 0 & -2 & -1 & -3 & 4 & -1 \\ 0 & 0 & -1 & -4 & -1 & 0 & 1 & 2 \end{bmatrix}$$

source

Pixel intensity (N^2 values)

DCT

Transformed coordinates (N^2 values)

How can we apply the quantizer?

Let's continue the example

We consider a quantizer matrix, in which higher frequency are quantized more heavily

Low frequencies

$$\begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

High frequencies

Let's continue the example

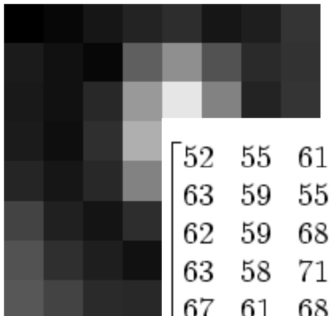

$$\begin{bmatrix} 52 & 55 & 61 & 66 & 70 & 61 & 64 & 73 \\ 63 & 59 & 55 & 90 & 109 & 85 & 69 & 72 \\ 62 & 59 & 68 & 113 & 144 & 104 & 66 & 73 \\ 63 & 58 & 71 & 122 & 154 & 106 & 70 & 69 \\ 67 & 61 & 68 & 104 & 126 & 88 & 68 & 70 \\ 79 & 65 & 60 & 70 & 77 & 68 & 58 & 75 \\ 85 & 71 & 64 & 59 & 55 & 61 & 65 & 83 \\ 87 & 79 & 69 & 68 & 65 & 76 & 78 & 94 \end{bmatrix}$$

source

$\rightarrow$ DCT $\rightarrow$

$$\begin{bmatrix} -415 & -30 & -61 & 27 & 56 & -20 & -2 & 0 \\ 4 & -22 & -61 & 10 & 13 & -7 & -9 & 5 \\ -47 & 7 & 77 & -25 & -29 & 10 & 5 & -6 \\ -49 & 12 & 34 & -15 & -10 & 6 & 2 & 2 \\ 12 & -7 & -13 & -4 & -2 & 2 & -3 & 3 \\ -8 & 3 & 2 & -6 & -2 & 1 & 4 & 2 \\ -1 & 0 & 0 & -2 & -1 & -3 & 4 & -1 \\ 0 & 0 & -1 & -4 & -1 & 0 & 1 & 2 \end{bmatrix}$$

DCT

Quantizer

Q matrix

$$\begin{bmatrix} 16 & 11 & 10 & 16 & 24 & 40 & 51 & 61 \\ 12 & 12 & 14 & 19 & 26 & 58 & 60 & 55 \\ 14 & 13 & 16 & 24 & 40 & 57 & 69 & 56 \\ 14 & 17 & 22 & 29 & 51 & 87 & 80 & 62 \\ 18 & 22 & 37 & 56 & 68 & 109 & 103 & 77 \\ 24 & 35 & 55 & 64 & 81 & 104 & 113 & 92 \\ 49 & 64 & 78 & 87 & 103 & 121 & 120 & 101 \\ 72 & 92 & 95 & 98 & 112 & 100 & 103 & 99 \end{bmatrix}$$

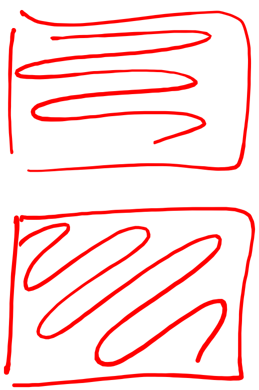

$$\begin{bmatrix} -26 & -3 & -6 & 2 & 2 & -1 & 0 & 0 \\ 0 & -2 & -4 & 1 & 1 & 0 & 0 & 0 \\ -3 & 1 & 5 & -1 & -1 & 0 & 0 & 0 \\ -4 & 1 & 2 & -1 & 0 & 0 & 0 & 0 \\ 1 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \\ 0 & 0 & 0 & 0 & 0 & 0 & 0 & 0 \end{bmatrix}$$

Quantized DCT

-26, -3, 0, -3, -2, -6,

→ (0, 30)

Can we reconstruct the image?

original

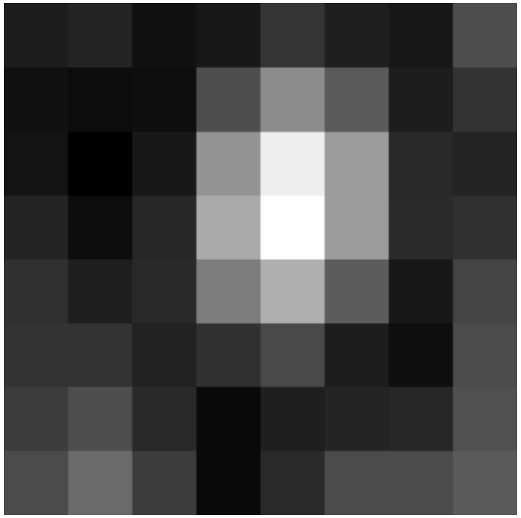

decompressed

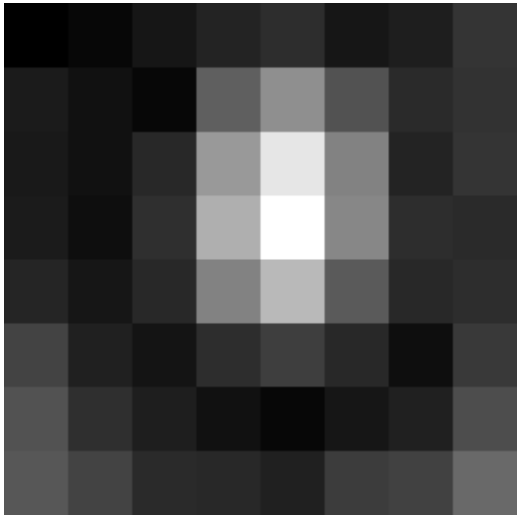

however visually the blocks are not very different

Hence, we highly compressed the image without having “perceptual” losses!

Source Coding

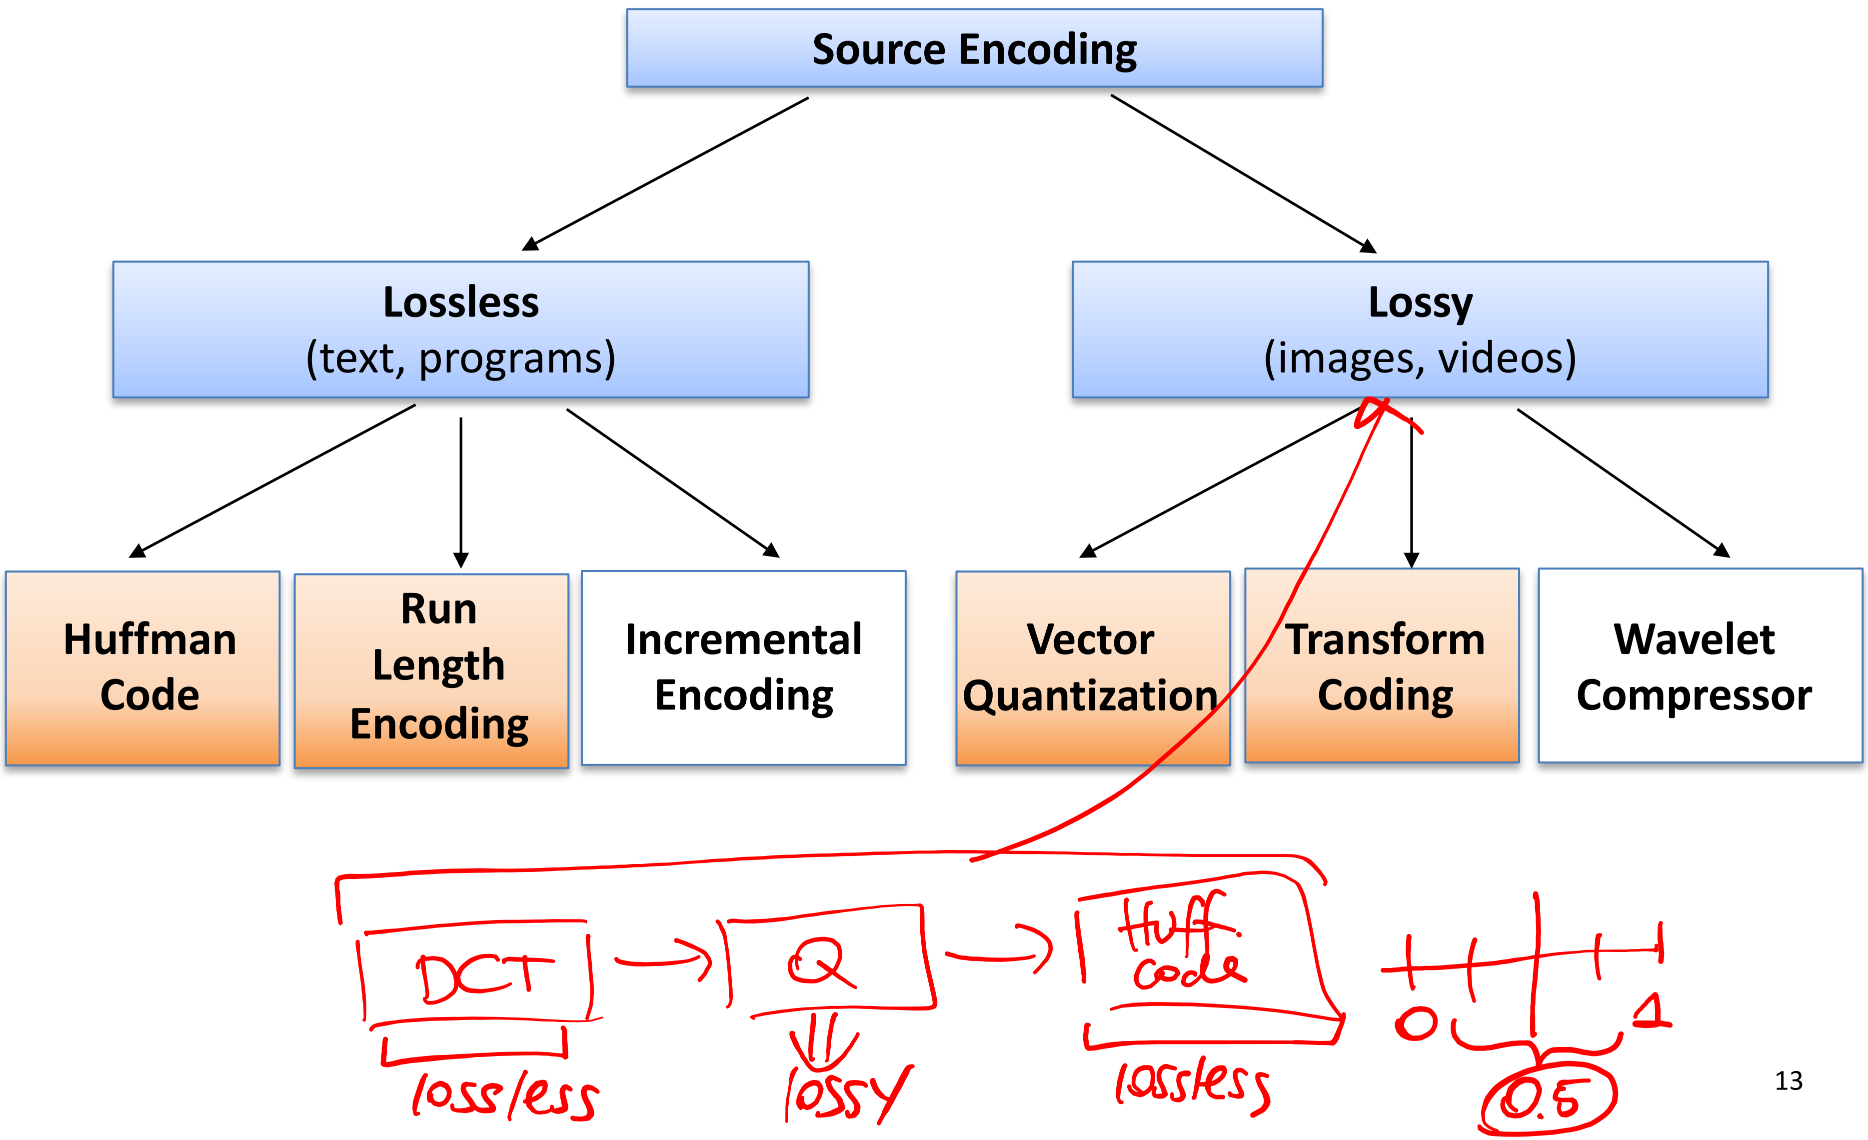

Take Home Message

- Information theory tells us to look at the probability of the source when coding (most probable words, shortest bits)
- Quantizer follows the same concept: quantize heavily the least probable ranges

- Image compression is composed of
  - There is redundancy in the images.
  - Artefacts in high-frequencies (edges) are less noticeable from the human eye
  - Source transformation (to compact the energy – most high frequency values will be low)
  - Quantizer (2D quantizer that will cut off high-frequencies)
  - Zig-zag rastering will maximize the probability of having 0s
  - Huffman lossless code as last step

Thank You